

# GEMM Benchmarks

deval-d@stanford.edu

June 6, 2026

From the last note it's clear that LAK does well for “shorter”  $C$  ( $m < 64$ ), but lags behind OpenBLAS for “taller”  $C$ . This is because LAK did not block  $C$  by MC, but worked on full contiguous column panels. As matrices got taller, full column-panels of  $C$  were not able to be kept close in cache for the next KC block of  $A$  and  $B$  to accumulate into.

I wrote a second GEMM implementation, meant for taller matrices that does block by MC. Internally, `gemm` uses the shorter matrix variant until  $m$  exceeds some preset constant threshold. Then it switches to using the taller matrix variant.

Below are the GEMM benchmarks for the common no-transpose  $\times$  no-transpose case. I made benchmarks to both highlight LAK's behavior for small  $m$  matrices and overall performance as  $n$  approaches 1792. I use the colorblind-safe [Okabe-ito](#) palette.

These benchmarks may change once I better tune blocking constants for the taller matrix implementation.

---

The benchmarks were run on Apple M4, 16GB unified memory.

All benchmarks are single-threaded and were run on square  $n \times n$  matrices. To compare against faer I set  $\beta = 0$  for the BLAS libraries, including `lak`, and  $\beta = \text{Accum::Replace}$  for faer. This makes the benchmark only time the operation  $C \leftarrow \alpha AB$ .

Each plot shows:

- LAK - (portable-simd, safe Rust)
- OpenBLAS armv8
- BLIS
- [faer](#)

# Small Matrices




Figure 1: Single-precision GEMM benchmark for small  $m$  matrices.




Figure 2: Double-precision GEMM benchmark for small  $m$  matrices.

# Complete Benchmark




Figure 3: Single-precision GEMM benchmark.




Figure 4: Double-precision GEMM benchmark.