

Lean Design of Optimized Level-1 BLAS

LAK Working Note #2

deval-d@stanford.edu

13 May 2026

Abstract

The level-1 BLAS consists of $O(n)$ vector-vector routines with low arithmetic intensity. In **lak**, these kernels fall into two implementation cases. Elementwise routines already expose independent work through ordinary scalar loops over contiguous slices, so LLVM can often generate SIMD code without manual vector intrinsics. Reduction routines are different: their natural scalar form introduces a loop-carried accumulator, which hides the available parallelism behind a serial recurrence. Those kernels need to be reshaped with independent partial sums and an explicit SIMD accumulator. This gives a lean level-1 design: keep the ordinary loops ordinary where the compiler can see the structure, and use manual SIMD only where the mathematical form of the routine would otherwise serialize the hot loop.

The compiler discussed is **rustc 1.94.0-nightly**. Other compilers may not have similar behavior.

1 Dependency Chains

The level-1 BLAS routines operate on vectors. Examples include scaling a vector, adding one vector to another, copying data, taking a dot product, and computing vector norms. These routines all have $O(n)$ work, but they do not all present the same optimization problem to the compiler. The important distinction in this note is whether the natural scalar loop has a *loop-carried dependency*. A loop-carried dependency occurs when one iteration of the loop needs a value produced by an earlier iteration. When no such dependency exists, the iterations can be treated as independent pieces of work. When such a dependency does exist, the loop looks sequential unless the implementation rewrites it to expose more parallelism.

The **AXPY** routine is the simple case. It computes

$$y_i \leftarrow \alpha x_i + y_i. \quad (1)$$

Each index i updates a different element of y . The calculation of y_7 does not depend on the result of the computation of y_6 . This is the elementwise structure shown on the left of Figure 1. Once the compiler sees a clean loop over contiguous slices, the elementwise structure is explicit. On a modern target, LLVM can often turn that scalar loop into SIMD instructions automatically, even though the source code never mentions SIMD directly. This will be covered more in sections 2 and 3.

The dot product has a different shape. It computes

$$x \cdot y = \sum_{i=0}^{n-1} x_i y_i. \quad (2)$$

A direct scalar implementation naturally introduces a running sum:

```
pub fn dot<T>(<
  x: VecRef<'_, T>,
  y: VecRef<'_, T>,
) -> T
where
  T: Default
    + Copy
    + Fma
    + Mul<Output=T>
    + Add<Output=T>,
{
  // ensures buffers have equal length
  assert_length_eq!(x, y);

  // access stored @[T] buffer;
  let xdata = x.as_slice();
  let ydata = y.as_slice();

  // running sum
  let mut sum = T::default();
  for (xi, yi) in xdata.iter().zip(ydata.iter()) {
    // dependency chain
    // sum += xi * yi
    sum = xi.fma(yi, sum);
  }

  sum
}
```

The inner loop does not compute an independent output for each index. Instead, each iteration reads the current value of **sum** and produces the next value of **sum**. The value produced at iteration i is therefore an input to iteration $i + 1$. This is the dependency chain shown on the right side of Figure 1.

This does not mean the dot product is impossible to vectorize. The products $x_i y_i$ are independent. What is serial in the naive loop is the single accumulated **sum** state. To use SIMD well, the implementation has to change the *shape* of the computation: instead of one running sum, it uses several independent partial sums and combines them at the end, outside of the inner loop. This rewrite is the main subject of section 3.



Figure 1: Classification of the level-1 kernels. Elementwise routines expose independent work directly. Dependency-chain routines carry state between iterations, hiding the available parallelism behind a scalar recurrence.

This gives the organizing principle for the rest of the note. For elementwise kernels, the main job is to expose a clean contiguous loop and let the compiler emit SIMD when it can. For reduction kernels, the main job is to remove the single scalar recurrence from the hot loop by introducing independent accumulators. The following two sections use **AXPY** and **DOT** as representative cases: **AXPY** shows when the clean scalar source is enough, while **DOT** shows when the computation itself has to be reshaped.

2 AXPY

$$y_i \leftarrow \alpha x_i + y_i \quad (3)$$

“Alpha x plus y,” or **AXPY**, performs a scaled vector addition. As discussed above, it is a purely elementwise operation. In **lak**, the contiguous-only vector interface gives the compiler a simple loop over two slices, with one independent update per index. In this memory-bound setting, that scalar source is enough for LLVM to generate a competitive SIMD loop. The implementation can therefore remain minimal:

```
pub fn axpy<T>(<
  a: T,
  x: VecRef<'_, T>,
  mut y: VecMut<'_, T>,
)
where
  T: Copy
    + AddAssign
    + Mul<Output=T>
    + Fma,
{
  assert_length_eq!(x, y);

  let x_slice = x.as_slice();
  let y_slice = y.as_slice_mut();

  // no simd needed, already fast
  for (xi, yi) in x_slice.iter().zip(y_slice.iter_mut()) {
    // *yi += a * xi + *yi
    *yi = a.fma(xi, *yi);
  }
}
```

***yi = a.fma(xi, *yi)** is the entire inner loop. There is no explicit vectorization here. The single-threaded, single-precision **SAXPY** benchmarks below test this implementation for vector lengths up to $2^{13} = 8192$.




Figure 2: Single-threaded **SAXPY** memory bandwidth benchmark relative to OpenBLAS armv8.




Figure 3: Single-threaded **SAXPY** time benchmark relative to OpenBLAS armv8. The y -axis is log-scaled.

In the bandwidth plot, higher is better. In the time plot, lower is better, and the y -axis is log-scaled. At $n = 8192$, **lak** finishes the routine in 754.7ns on average. OpenBLAS armv8 takes 1.541 μ s on average on the same **LB85.7** benchmark.

The compiler-generated assembly for single precision shows that the scalar source loop was automatically vectorized into a modestly unrolled SIMD loop:

```
LB85_7:
ldp q1, q2, [x9, #-32]
ldp q3, q4, [x9], #64
ldp q5, q6, [x10, #-32]
ldp q7, q16, [x10]
fmla.4s v5, v0, v1
fmla.4s v6, v0, v2
fmla.4s v7, v0, v3
fmla.4s v16, v0, v4
stp q5, q6, [x10, #-32]
stp q7, q16, [x10], #64
subs x11, x11, #16
b.ne LB85_7
cmp x1, x8
b.eq LB85_15
tst x1, #0xc
b.eq LB85_13
```

Each loop iteration loads 16 f32 values from x into four SIMD registers, **q1** through **q4**. It also loads 16 f32 values from y into **q5**, **q6**, **q7**, and **q16**. The core operation is:

```
fmla.4s v5, v0, v1
fmla.4s v6, v0, v2
fmla.4s v7, v0, v3
fmla.4s v16, v0, v4
```

v0 is alpha broadcast across all four lanes. The updated 16 floats are then stored back into y . The loop is not scalar, but it is also not aggressively unrolled. For a memory-bound **AXPY**, that is the right outcome: the compiler has widened the loop enough to make arithmetic throughput a secondary concern, while the limiting cost remains the load and store traffic. Figure 4 shows this source-level independence being widened into SIMD work by the compiler.




Figure 4: Independent elementwise iterations are automatically widened into a single SIMD operation.

3 DOT

The dot product

$$\vec{x} \cdot \vec{y}, \quad (4)$$

has a different implementation shape because its direct scalar form carries a running sum.

```
pub fn dot<T>(<
  x: VecRef<'_, T>,
  y: VecRef<'_, T>,
) -> T
where
  T: Default
    + Copy
    + Fma
    + Mul<Output=T>
    + Add<Output=T>,
{
  assert_length_eq!(x, y);

  let xdata = x.as_slice();
  let ydata = y.as_slice();

  // running sum
  let mut sum = T::default();
  for (xi, yi) in xdata.iter().zip(ydata.iter()) {
    // dependency chain
    // sum += xi * yi
    sum = xi.fma(yi, sum);
  }

  sum
}
```

The generated assembly for this routine shows the following inner loop:

```
LB85_1:
ldr s1, [x0], #4
ldr s2, [x2], #4
fmadd s0, s1, s2, s0
subs x8, x8, #1
b.ne LB85_1
```

This loop uses **fmadd**, a scalar f32 FMA. It is not the **fmla** instruction shown in the **AXPY** assembly above, which is a vector FMA. The loop performs one scalar FMA per iteration. The accumulator lives in **s0**, and every iteration reads and writes that same register. LLVM therefore preserves the scalar recurrence rather than widening the loop into SIMD instructions as it did for **AXPY**. In this form, the routine takes about 6 μ s for vectors of length $n = 8192$, while OpenBLAS armv8 takes about 1 μ s on the same benchmark.

The issue is not that the products $x_i y_i$ are serial. They are independent. The issue is that the source code asks the machine to route all of those products through one accumulated scalar state. The fix is to replace the single long recurrence with several independent partial recurrences. Each SIMD lane carries its own partial sum through the hot loop, and those partial sums are reduced only after the main loop is finished. This requires manual SIMD.

```
pub(crate) const LANES: usize = 32;

pub fn dot<T>(<
  x: VecRef<'_, T>,
  y: VecRef<'_, T>,
) -> T
where
  T: SimdElement
    + Copy
    + Default
    + AddAssign
    + Mul<Output=T>
    + Add<Output=T>
    + Fma,
{
  Simd<T, LANES>: SimdFloat<Scalar=T>
    + AddAssign
    + Fma,
{
  assert_length_eq!(x, y);

  let x_slice = x.as_slice();
  let y_slice = y.as_slice();

  // partitioning buffers into fixed LANES lengths
  // with a leftover tail
  let (x_chunks, x_tail) = x_slice.as_chunks::<<LANES>();
  let (y_chunks, y_tail) = y_slice.as_chunks::<<LANES>();

  // inner loop
  // independent accumulators in each LANE
  let mut accumulator = Simd::<T, LANES>::splat(T::default());
  for (ix_chunk, iy_chunk) in x_chunks.iter().zip(y_chunks.iter()) {
    let x_vec = Simd::from_array(ix_chunk);
    let y_vec = Simd::from_array(iy_chunk);

    accumulator = x_vec.fma(y_vec, accumulator);
  }

  // scalar tail leftover
  let mut sum = T::default();
  for (ix_t, iy_t) in x_tail.iter().zip(y_tail.iter()) {
    sum = xt.fma(yt, sum);
  }

  // final horizontal sum across all LANES
  // outside the inner loop
  accumulator.reduce_sum() + sum
}
```

The generics get a little ugly, but the idea is simple. A chunk size **LANES** is fixed at compile time. The vectors x and y are split into full chunks of length **LANES**, plus a tail of length less than **LANES**. In the main loop, each chunk is loaded into a SIMD vector, and the SIMD **fma** updates the accumulator vector. Each lane of that accumulator carries its own partial sum. I found **LANES=32** to be fastest for this benchmark.

After the SIMD loop, the leftover tail is handled with a scalar loop. Then the SIMD accumulator is horizontally reduced and added to the scalar tail sum to produce the final dot product $\vec{x} \cdot \vec{y}$. The hot loop has been changed from one scalar recurrence into several independent recurrences. The parallelism was already present in the dot product, but the naive implementation hid it behind a single running sum. This process is visualized in Figure 7. The generated assembly for the improved DOT inner loop is shown below:

```
LB85_4:
ldp q17, q16, [x9, #64]
ldp q19, q18, [x9, #32]
ldp q20, q21, [x9]
ldp q23, q22, [x10, #32]
ldp q24, q25, [x10]
fmla.4s v7, v25, v21
fmla.4s v6, v24, v20
fmla.4s v5, v23, v19
ldp q20, q19, [x10, #64]
fmla.4s v4, v22, v18
fmla.4s v3, v20, v17
fmla.4s v2, v19, v16
ldp q17, q16, [x9, #96]
fmla.4s v1, v19, v17
fmla.4s v0, v18, v16
add x10, x10, #128
add x9, x9, #128
subs x8, x8, #1
b.ne LB85_4
```

This is the manually vectorized dot-product loop. Unlike the naive implementation, it keeps eight independent SIMD accumulators, **v0** through **v7**. Each accumulator carries a different partial sum. One loop iteration processes 32 elements of x and 32 elements of y : eight SIMD registers, with four f32 lanes per register. The loop therefore performs eight **fmla.4s** instructions per iteration, instead of one scalar **fmadd** on one pair of elements.

The important point is not only that the improved loop uses SIMD. The important point is that it changes the dependency structure. There are still dependencies, because each accumulator is carried across loop iterations. But there are now eight independent vector accumulation chains instead of one scalar running sum. That gives the CPU more independent work to schedule while FMA results are still waiting on latency.

Figure 5: Single-threaded **SDOT** memory-bandwidth benchmark relative to OpenBLAS armv8.

Figure 6: Single-threaded **SDOT** time benchmark relative to OpenBLAS armv8.

Figure 7: A running accumulator creates a loop-carried dependency. Splitting the work across independent accumulators exposes parallelism inside the loop; the partial sums are reduced afterward.

The level-1 BLAS does not need every routine to be written as a hand-vectorized SIMD kernel. For elementwise routines like **AXPY**, the scalar source already exposes independent work, and LLVM can usually generate the SIMD loop on its own. For dependency-chain routines like **DOT**, the scalar source hides the available parallelism behind a single recurrence. Those routines need a different implementation shape: independent partial sums in the hot loop, followed by a final reduction. This is the main design principle used in **lak**'s level-1 kernels. Keep the ordinary loops ordinary where the compiler can already see the structure. Use manual SIMD only where the mathematical structure of the routine would otherwise serialize the computation.